

Generative Adversarial Networks(GANs)

Anuj Bansal (2018201096)

Nilabja Bhattacharya (2018201036)

Why Generative Models?

- **In discriminative models**

- Given an image X , predict a label Y
- Estimates $P(Y|X)$

- **Discriminative models have several key limitations**

- Can't model $P(X)$, i.e. the probability of seeing a certain image
- Thus, can't sample from $P(X)$, i.e. can't generate new images

- **Generative models (in general) cope with all of above**

- Can model $P(X)$
- Can generate new images

Magic of GAN

Which one is Computer generated?



Adversarial training

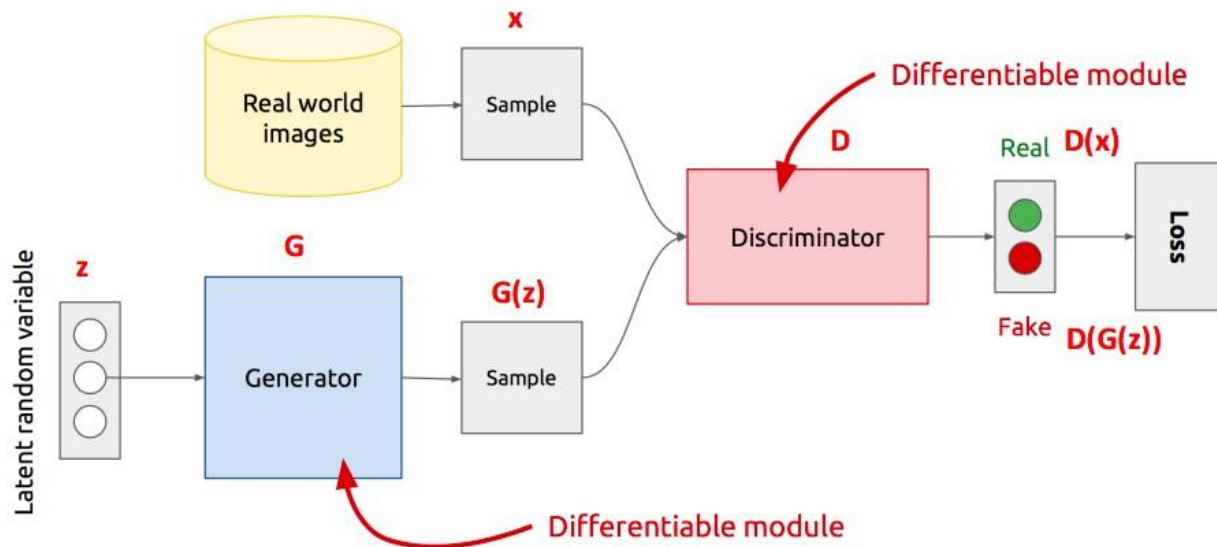
- **In discriminative models**

- We can generate adversarial samples to fool a discriminative model
- We can use those adversarial samples to make models robust
- We then require more effort to generate adversarial samples
- Repeat this and we get better discriminative model

- **GANs extend that idea to generative models**

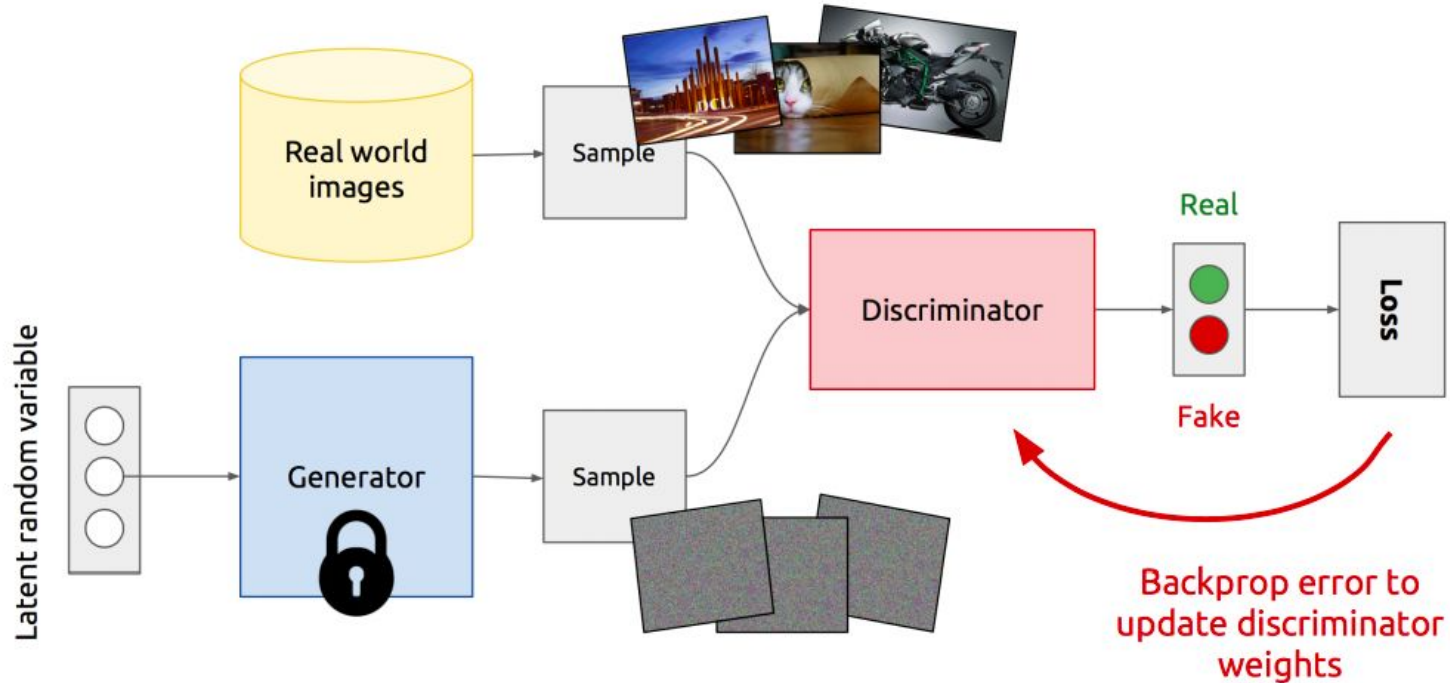
- Generator: generate fake samples, tries to fool the Discriminator
- Discriminator: tries to distinguish between real and fake samples
- Train them against each other
- Repeat this and we get better Generator and Discriminator

GAN's Architecture

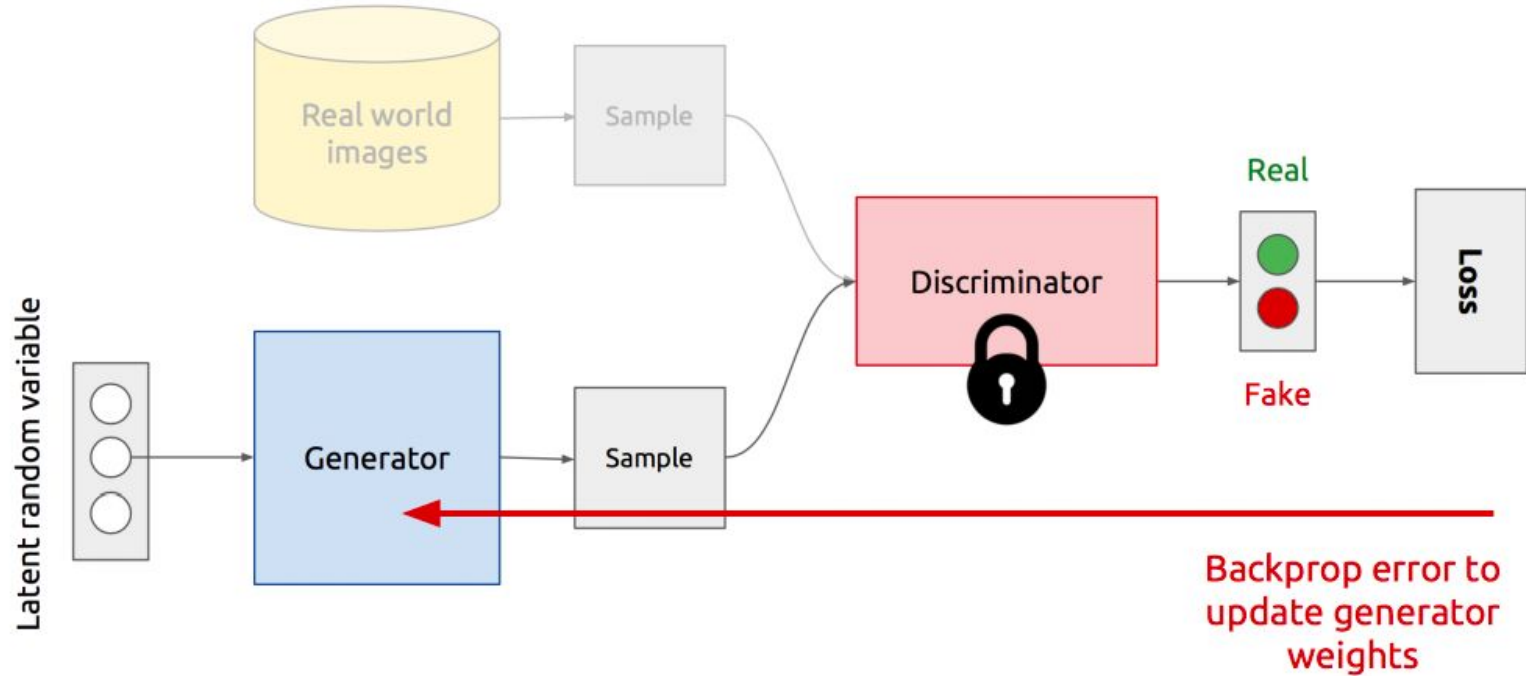


- **z** is some random noise (Gaussian/Uniform).
- **z** can be thought as the latent representation of the image.

Training Discriminator



Training Generator



GAN's formulation

$$\min_G \max_D V(D, G)$$

- It is formulated as a **minimax game**, where:
 - The Discriminator is trying to maximize its reward $V(D, G)$
 - The Generator is trying to minimize Discriminator's reward (or maximize its loss)

$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

- The Nash equilibrium of this particular game is achieved at:
 - $P_{data}(x) = P_{gen}(x) \quad \forall x$
 - $D(x) = \frac{1}{2} \quad \forall x$

GAN's Algorithm

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Discriminator
updates

Generator
updates


Advantages of GAN

- **Plenty of existing work on Deep Generative Models**
 - Boltzmann Machine
 - Deep Belief Nets
 - Variational Auto Encoders (VAE)
- **Why GANs?**
 - Sampling (or generation) is straightforward
 - Training doesn't involve Maximum Likelihood estimation
 - Robust to Overfitting since Generator never sees the training data
 - Empirically, GANs are good at capturing the modes of the distribution


Problems with GAN

- **Probability Distribution is Implicit**
 - Not straightforward to compute $P(X)$
 - Thus Vanilla GANs are only good for Sampling/Generation.
- **Training is Hard**
 - Non-Convergence
 - Mode-Collapse

Result of GAN



Thank you



—